

# Stats Seminar (STATS 770)

## McMaster University

### R programming assignment

15 November 2016

1. Consider a random sample from a mixture distribution with probability density function

$$f(y; \mu, \sigma) = \pi f_1(y; \mu, \sigma) + (1 - \pi) f_2(y; \mu, \sigma), \quad y \in \mathbb{R}$$

where  $f_1(\cdot)$  and  $f_2(\cdot)$  are the PDFs of the Normal distribution with  $(\mu=1.5, \sigma=1)$  and  $(\mu=2, \sigma=1.75)$  respectively.

$$\pi = \begin{cases} 1 & \text{with prob. } P(U \geq 0.5) \\ 0 & \text{with prob. } P(U < 0.5) \end{cases}$$

where  $U \sim \text{Uniform}(0,1)$ .

- Create a numeric vector containing  $10^4$  values of  $f(y; \mu, \sigma)$ .
  - Calculate the mean and standard deviation of the sample using base R functions.
  - Calculate the mean and standard deviation of the sample explicitly using a `for` loop to add each term (without using default R functions).
  - Compare the computation time using `system.time()` and report for both the cases above (i.e. with and without default R functions).
  - Compare the computation time using a function from a benchmarking package (e.g. `microbenchmark::microbenchmark` or `rbenchmark::benchmark`) and report for both cases.
  - Using `Rcpp::cppFunction`, compute the mean and standard deviation of the sample and report the computation time using a benchmark package. *You may use either explicit for loops or Rcpp “sugar” functions* (e.g. see here or particularly here).
2. A standard test case for simple computational methods is to estimate the value of  $\pi$  by the rejection method. In other words:
    - Pick  $N$  values uniformly in the unit square  $[0, 1] \times [0, 1]$ .
    - Count the fraction that fall inside the unit circle (i.e.  $x^2 + y^2 < 1$ ) (in R you can conveniently do this via `mean(x^2+y^2<1)`, because applying `mean()` converts logical variables (FALSE/TRUE) to numeric (0/1)).
    - This value should be the area of the quarter-circle divided by the area of the square, i.e.  $(\pi r^2/4)/r^2 = \pi/4$ ; multiplying by 4 gives  $\pi$ .
  - Write an R function `approx_pi` (with argument `N` specifying the number of samples) to compute this approximation. Try it out.
  - Now write R code to do this in parallel across more than one core on your machine. On Linux or MacOS `library(parallel); mclapply(c(N/2, N/2), approx_pi)` should work. If you are on Windows you will have to do something like

```
library(parallel)
cl <- makeCluster(2)
parLapply(cl, c(N/2, N/2), approx_pi)
stopCluster(cl)
```

Use `system.time()` and/or a benchmarking package to compare the performance of the parallel and non-parallel codes.